c

Multi-agent execution (each contribution measured)

no LLM hard failures only *benefit falsified (F19/P3)* *19/19 with 0 invalid* *ClaimGate-gated*

governance: an LLM advisor can never flip a successful run to failed
(review_unavailable → zero-confidence pass, error retained for audit)