Three failure modes, all measured in this work

1 Uncalibrated tool trust

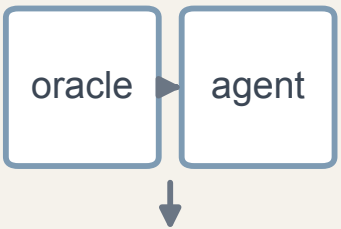

score consumed at face value

no power check for task, family or regime

2 Claim overreach

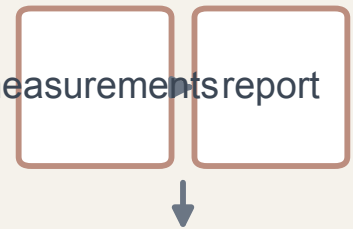

claim exceeds the evidence

overreach enters through language

3 Confounded improvement

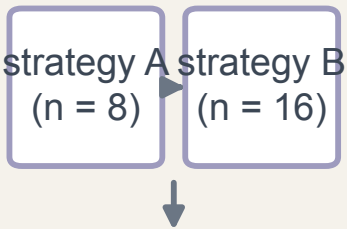

apparent gain = more samples

budget confound, not mechanism